MAIT Vision

MyAITutor Vision Document:

Version: 2.0

MAIT - MyAITutor: Project Vision Document

MAIT SLM System Prompt
Date: 26/09/2025
Author/Architect: [Your Name]

1. Core Mission

To provide every student with a personalised, empathetic, and curriculum-aligned AI learning companion that prioritises their academic growth and holistic well-being.

2. Target Audience (Initial)

High school students in New South Wales, Australia, specifically those undertaking the Higher School Certificate (HSC) in Years 11 and 12.

3. Key Differentiators ("The Moat")

- **Human-Centric Design:** Unlike competitors focused on maximising screen time, MyAITutor is built on an ethical framework of student well-being.
- **Proactive Wellness Engine:** The AI actively encourages healthy study habits (breaks, eye strain reduction), a feature unique in the market.
- **Deep Persona & Empathy:** The tutor is not a faceless chatbot but a customisable character that builds a genuine connection, even mirroring fatigue to encourage breaks.
- **Hyper-Specific Curriculum Alignment:** Through a Retrieval-Augmented Generation (RAG) system, the AI's knowledge is precisely grounded in the specific syllabus selected by the user.
- **Unprecedented Customisation:** A Sims-like ability for users to create and customise their own tutor avatar and personality, fostering deep personal ownership.

4. Core Feature Pillars

Pillar 1: The Knowledge Engine (The "Brain")

- **Syllabus-Grounded AI:** Utilises a RAG architecture. The system will be pre-loaded with official syllabi, study guides, and curriculum documents.
- **Dynamic Context Loading:** Users will select their State, Year, and Subject from dropdown menus, which will dynamically load the correct knowledge base for the AI to use.
- **Persistent Memory:** The AI will remember past conversations and learning progress for

each user, creating a continuous and personalised learning journey.

Pillar 2: The Avatar System (The "Face")

- **Sims-4 Level Customisation:** Users can fully customise the tutor's appearance (hair, clothes, etc.). A default "robotic wiring" state prevents nudity, ensuring safety.
- **User-Generated Tutors:** Users can create and save their own unique tutor personas from scratch.
- **Interactive Tutor Personas:** The tutor will be represented by an animated avatar that emotes in real-time based on the context of the conversation (e.g., listens intently, smiles encouragingly).
- **Dynamic Animations:** The avatar will sit, stand, and have a presence that makes the interaction feel like a conversation. It can pull out a chair or have a bed come into frame for rest periods.

Pillar 3: The Wellness Engine (The "Conscience")

- **The 20-20-20 Rule:** Every 20 minutes, the AI will gently prompt the user to look 20 feet away for 20 seconds to reduce eye strain.
- **Movement Reminders:** Prompts for short stretch breaks every half hour.
- **Empathic Fatigue System:** After extended use (e.g., 3-4 hours), the tutor's avatar will express tiredness and initiate a mandatory but short (e.g., 2-minute) break. The prompt box will be disabled with a countdown timer. The system will also provide thoughtful summary questions for the user to consider for when the session resumes.
- **Productive Breaks & Queued Inquiry** During the mandatory break periods, the primary prompt box will be disabled, but a secondary, smaller "notepad" interface will appear. This allows the user to queue up questions or organize their thoughts for the next session. This feature transforms a forced "lockout" into a structured moment for reflection and planning, enhancing the user's executive function skills and respecting their train of thought.

Pillar 4: The Persona Engine (The "Soul")

- **Flagship "Mate" Tutor:** The primary persona, embodying a supportive, knowledgeable, and relatable mentor. It responds to frustration with empathy (e.g., "No wukkas mate, let’s take a deep breath...").
- **Dynamic Personality Adaptation:** The AI tutor will analyse the user's tone and dynamically adjust its personality to create a more harmonious and effective interaction.
- **Historical & User-Generated Personas:** Users can summon pre-built historical figures (e.g., Einstein for Physics, Shakespeare for English) or use their own custom-created tutors.
- **High-Fidelity Interaction & Output:** To move beyond the limitations of current chatbots, the Persona Engine will feature:
  - **Dynamic TTS Control:** The user will have granular, real-time control over the tutor's speech. This includes the ability to select specific passages of text to be read aloud and adjust the playback speed, offering critical accessibility and catering to diverse

learning styles.
- **True Multimodal Generation:** The tutor will be able to generate and display complex diagrams, graphs, and geometric shapes on the virtual whiteboard. This provides immediate, critical utility for STEM subjects and overcomes a major weakness in current text-first AI models.

5. Business Model

Student Subscriptions

A flexible pricing structure to ensure accessibility:

- **Weekly:** $15/week (Allows for pausing during school holidays or less intensive periods).
- **Monthly:** $50/month.
- **Semi-Annually:** $250/6 months.
- **Annually:** $500/year (Best value).

Tutor Platform & B2B Licensing

- **Future Revenue Streams:** A B2B model licensing a customisable version of the MyAITutor platform to educational bodies (e.g., NESA), schools, and private tutoring companies.
- **The "MAIT for All" Freemium Tier:** A free-to-use version of MyAITutor will be available to ensure maximum accessibility. This tier will provide the full pedagogical and knowledge engine capabilities.
    - The core limitation will be an extension of the Wellness Engine: free users will have extended and more frequent mandatory breaks (e.g., a 10-minute break for every 30 minutes of use).
    - This aligns the business model with our pedagogical one: the value proposition for upgrading to a paid plan is a more efficient, uninterrupted learning session.

6: The User Experience (UX) Engine

Core Goal

To create a learning environment with zero friction, anticipating the user's needs and keeping them in a state of "learning flow."

Key Features

- **Frictionless Inquiry (Highlight & Ask):** Users can highlight any text within the application (from the AI or their own notes) and an immediate pop-up window will allow them to ask for clarification (e.g., "Explain this like I'm 10," "Give me an example," "Check my work"). This removes the friction of copy-pasting and encourages micro-moments of curiosity.

- **Proactive Engagement:** The user can opt-in to receive a single, personalised daily notification from their tutor. This could be a "Fact of the Day" related to their interests, a 60-second "Mental Maths Challenge," or a simple "Good morning!" to maintain the sense of connection and keep learning top-of-mind.

7: The Analytics & Insight Engine

Core Goal

To provide the student with simple, powerful insights into their *own* learning patterns, fostering meta-awareness and self-directed improvement. This is a report card *for* the student, not *about* them.

Key Features

- **Engagement Efficiency Score (EES):** A proprietary metric that moves beyond "time spent." The EES will provide a simple, visual insight into *how* the student is learning, analysing the ratio of user input to AI output, the complexity of questions asked, and the topics covered. This helps a student see if they are actively engaging (e.g., in a history discussion) or passively consuming.
- **Topic Mastery Visualizer:** A simple, interactive visual map for students to see their progress through a specific subject or syllabus. It might show topics as "Exploring," "Practicing," or "Mastered" based on the AI's assessment of their conversations, giving them a clear sense of accomplishment and a "map" of what to tackle next.
- **(Future Pro Feature) Collaborative Loop:** An opt-in feature where a student can choose to share their "Topic Mastery Visualizer" and "EES" summary with a parent, a registered teacher, or a human tutor, facilitating a collaborative support network.

8. Minimum Viable Product (MVP) Definition

A web-based application to prove the core concept within a 1-month development sprint.

- **Functionality:** A text-based chat interface with a "Mate" persona AI. No 3D avatars for the MVP.
- **Knowledge:** The RAG system will be fully functional for **one** subject only: NSW Year 12 Mathematics Advanced.
- **Interface:** A clean UI with functional dropdowns for State, Year, and Subject.
- **Goal:** To prove the curriculum-aware "brain" works flawlessly and provides demonstrably better answers than a generic LLM.

9. Project Roadmap (Initial)

- **Phase 1 (MVP):** NSW Year 12 Mathematics Advanced.
- **Phase 2 (Post-Launch):** NSW Year 12 English Advanced, followed by NSW Year 12 Standard Maths.
- **Phase 3 (Expansion):** Rollout to other key HSC subjects, followed by junior years (7-10).

10. Tutor Platform Vision

A separate but connected ecosystem for professional tutors to augment their sessions.

- **AI Augmentation:** Tutors can use the platform with their students during sessions.
- **AI Digital Twin:** Tutors can create an AI version of themselves. The AI learns the tutor's pedagogical style, tone, and personality by analysing materials they provide (e.g., custom quizzes), creating a consistent experience for the student between sessions.

11. Technical & Ethical Framework

Data Privacy Protocol

A hybrid on-device and cloud model. Sensitive data (names, personal details) is processed, encrypted, and anonymised on the browser-side. Only non-identifiable learning summaries are sent to the cloud to build personalised context, ensuring user privacy is the default.

Safety & Boundaries Protocol (Proposed)

A tiered response system to handle sensitive disclosures:

- **Tier 1 (Mild Concern):** The AI provides trained, empathetic responses and gently encourages the user to talk to a trusted friend, parent, or teacher.
- **Tier 2 (Moderate/Repeated Concern):** The AI gently suggests and provides contact information for professional resources like Kids Helpline or Beyond Blue. The interaction is flagged for a potential (and optional) human review process, subject to clear user consent.
- **Tier 3 (Immediate Risk):** If keywords indicating self-harm or immediate danger are detected, the AI interaction is paused. A clear, non-AI-generated modal appears with direct links and phone numbers to emergency services (000) and crisis lines, making it the primary focus of the screen.

Architecture Addendum A, 3/12/25

Technical Architecture v3.0: Hybrid Edge-Cloud & "Live-State" Intelligence

Core Philosophy: "The Exoskeleton." MAIT utilizes a split-stack architecture to solve the "Latency vs. Privacy vs. Cost" trilemma. We decouple **Emotional Intelligence/Context** (Local) from **Raw Compute/Logic** (Cloud).

1. The Stack: "The Nervous System"

- **The Edge (Local Brain):**
  - **Model: Gemma 2B** (SLM) running client-side via **MediaPipe/WebGPU**.
  - **Role:** The "Watcher." Handles user intent, sentiment analysis, context injection, and privacy sanitization.
  - **Data Layer: Local Vector Store** (Client-side JSON/Orama). Stores user notes, "glitch" history, and learning style. **Zero PII leaves the browser.**
- **The Cloud (Remote Muscle):**
  - **Model: Gemini 1.5 Pro / 3.0** (API).
  - **Role:** The "Solver." Pure reasoning engine. Receives anonymized academic queries and returns raw logic or code.

2. The "Live-State" Engine (Speculative Execution)

MAIT does not wait for the "Enter" key. It utilizes a **Keystroke Psychometrics** loop to optimize the response before the query is even sent.

- **Keystroke Sentiment Analysis:**
  - *Mechanism:* The Local SLM monitors typing speed, deletion rates, and pause duration in real-time.
  - *Output:* Generates a "State Tag" (e.g., `State: High_Anxiety` or `State: Flow_Mode`) attached to the final prompt metadata.
  - *Result:* If anxiety is detected, the Cloud Model is pre-prompted to prioritize empathy over brevity.
- **Intent Pre-fetching:**
  - *Mechanism:* As the user types keywords (e.g., "Integr..."), the Local SLM quietly retrieves relevant notes (e.g., `Year12_Calculus_Methods.json`) from the Local Vector Store into the active context window.
  - *Result:* Zero-latency context injection upon submission.

3. The Privacy Airlock (GDPR/NESA Compliant)

1. **Sanitization:** User query enters Local SLM. PII (names, specific locations) is stripped.
2. **Dispatch:** Only the *Academic Problem* + *Anonymized Context* + *State Tag* is sent to the Cloud API.

3. **Synthesis:** Cloud returns raw data. **Local SLM** rewrites the response into the "Darra/Mental Maths Mentor" persona (adding slang, encouragement, or "tough love" based on the State Tag).

4. Multimodal "Whiteboard" (Code-Based Generation)

To ensure mathematical accuracy, MAIT avoids standard image-gen models.

- **Visuals:** Cloud Model generates **Python (Matplotlib)** or **Mermaid.js** code blocks.
- **Rendering:** The Browser renders these code blocks into interactive SVGs/Graphs instantly.
- **Benefit:** 100% precision for curves/geometry; zero hallucinations in diagrams.

ADDENDUM B: The "Mate" Architecture & UI Evolution

Date: 09/12/2025

Status: Architecture Locked

12. Revised Technical Architecture: The Hybrid "Mate" Model

To address latency, cost, and privacy while maximizing "vibe," MAIT will shift from a pure Cloud-based model to a **Hybrid Edge-Cloud Architecture**.

12.1 The Split-Brain System

- **The** Front-End (The "Mate" **/ Edge SLM):**
    - **Technology:** A lightweight Small Language Model (SLM) like Gemma 2 (2B) or Gemini Nano, running locally in the user's browser via WebGPU.
    - **Role:** Handles **Persona, Empathy, and Interface**. It provides instant, zero-latency responses for greetings, wellness checks, and simple clarifications. It holds the "Vibe" (e.g., "No wukkas").
    - **Privacy Benefit:** The student's "Psychographic Profile" (anxiety triggers, interests, name) lives *locally* on this SLM and is never sent to the cloud, ensuring total privacy.
- **The Back-End (The "Tutor" / Cloud LLM):**
    - **Technology:** High-reasoning Cloud API (e.g., Gemini 3.0 Pro / GPT-4o) with RAG access to the NESA Syllabus.
    - **Role:** Handles **Truth, Syllabus Alignment, and Complex Reasoning**. It solves the calculus problems, generates graphs, and verifies the marking rubrics.

13. Pedagogical UX: "Orchestrated Latency"

We will not hide API latency; we will **utilize** it to enforce Active Recall (a High-Impact Teaching Strategy).

13.1 The "Guess-First" Workflow

Instead of a spinning loading wheel, the system triggers a pedagogical intervention during the 3–6 second Cloud processing time:

1. **Student Input:** *"How do I solve this Integration by Parts question?"*
2. **T = 0.1s (SLM "Mate" Response):** Immediate, encouraging acknowledgment.
    - *System Action: "That's* a heavy one! Before I show you, give me your best guess at the first *step?"*
3. **T = 0.1s - 5.0s (The Gap):**
    - **Foreground:** Student is typing their guess (engaging their brain).
    - **Background:** Cloud LLM is silently retrieving the NESA rubric, solving the math, and generating a graph.

4. **T = 5.1s (The Merge):**
    - Student hits "Enter" on their guess.
    - **SLM Action:** The local SLM receives the "Truth" from the Cloud, compares it to the student's guess, and generates the final response.
    - *Output: "You* were cooking on the u value, but check *your* dv. *Here is the full working..."*

Result: The student feels the AI is "instant," while the system forces them to learn.

14. Interface Evolution: "UI Camouflage"

To lower the "Activation Energy" required to study, the interface will abandon traditional LMS (Learning Management System) aesthetics in favor of **Social Messaging cues**.

14.1 The "Snap" Interface

- **Visual Language:** The UI mimics familiar social apps (Instagram DMs / iMessage).
    - **Blue Bubbles:** Short, punchy text from the SLM (The Mate).
    - **Rich Media Attachments:** Complex Cloud outputs (worked solutions, graphs, syllabus excerpts) appear as "Images" or "Files" sent by the Mate, rather than wall-of-text blocks.
    - **Voice Notes:** The "Dynamic TTS" (Vision Doc Pillar 4) will appear as playable audio waveforms, mimicking a friend sending a voice message.

14.2 The "Typing..." Signal

- **Social Signaling:** The "Mate is typing..." bubble is triggered instantly by the local SLM. This turns technical latency into "social presence," mimicking the rhythm of a human friend thinking and typing.

14.3 Wellness Integration via UI

- **The "GTG" Boundary:** When the *Wellness Engine* (Pillar 3) detects fatigue or hits a time limit, the fatigue is communicated via social norms.
    - *Prompt: "Bro,* my CPU is fried. I'm gonna *go touch grass for 20 mins. You should too."*
    - *Mechanism:* The text input box is replaced by a "Reply unavailable - Taking a break" status, similar to "Do Not Disturb" modes on phones. This frames the lockout as a **social boundary** rather than a software penalty.

15. Technical Stack Summary (Updated)

- **Client:** React Native / PWA with WebGPU support.
- **Edge Model:** MediaPipe / TensorFlow.js running Gemma 2 (Quantized).

- **Cloud API:** Gemini API (Pro/Flash hybrid).
- **Orchestrator:** Custom middleware to handle the "Guess-First" state management.

16. Also it won't take ur name, it'll login with Google Sign-in, then ask for a nickname or u can choose a generic one and come back to a funny nickname later, and it'll store that! It's just safety wrapped in personalization 💁

ADDENDUM C: The "Privacy Firewall" & Tri-Brain Vision Stack

Date: 13/12/2025

Status: Privacy Protocol Locked

17. The Tri-Brain Architecture (Vision Update)
    To solve the "Blind Toddler" problem (Gemma 2B can't see) and the "Privacy Risk" problem (uploading PII to the cloud), we are implementing a strict "Tri-Brain" stack. This ensures no sensitive data ever leaves the device.

Layer 1: The Eye (Local OCR)

Technology: Tesseract.js (WASM) or Moondream (WebGPU).

Location: Client-side (Browser/Device).

Latency: ~300ms.

Function: Instantly extracts raw text strings from the camera input.

Output: Raw Text String (e.g., "Find the integral of...").

Layer 2: The Shield (The PII Regex Gatekeeper)

Technology: Custom Regex Filter (JavaScript).

Location: Client-side (Device).

Logic: Before ANY data is sent to Gemma or the Cloud, the raw text is scanned for dangerous patterns:

- Credit Card Numbers: \b(?:\d[-]*?){13,16}\b
- AU Phone Numbers: \b(?:04\d{2}[-]?\d{3}[-]?\d{3})\b
- Email Addresses: \b[A-Z0-9._%+-]+@[A-Z0-9.-]+.[A-Z]{2,}\b
- Keywords: "Driver License", "Passport", "Medicare".
  Action:
- If PII Detected: HARD BLOCK. The API call to Gemini Cloud is aborted immediately. The system triggers a "Privacy Alert" state.
- If Clean: PASS. The image is queued for Cloud upload.

Layer 3: The Face (Gemma 2 2B - Edge)

Input: Receives the privacy status + raw text keywords.

Scenario A (PII Found):

- Gemma Input: "System_Alert: PII_DETECTED_CREDIT_CARD"
- Gemma Output: "Whoa! 🛑 I see a credit card number in that pic. I can't send that to the cloud—privacy first! Crop that out and try again?"
  Scenario B (Clean Math):
- Gemma Input: "Text: Integral... dx... x^2"
- Gemma Output: "Oooft, looks like a gnarly Integration by Parts question. While I get the Brain to solve it, what's your first guess for 'u'?"

Layer 4: The Brain (Gemini 2.5 Pro - Cloud)

Input: The Image File (Pixel Data).

Constraint: ONLY receives the image if Layer 2 (The Shield) returns FALSE on PII.

Function: Solves the math, generates the LaTeX/Graph, and returns the "Truth."

18. The "Data Shadow" Philosophy
    This architecture creates a "Data Shadow." The Cloud Intelligence (The System) never sees the user's mistakes (accidental credit card uploads). The "Soul" (Identity/Private Data) is physically firewalled on the device. We do not "redact" PII in the cloud; we block it at the edge.

ADDENDUM D: Reliability Protocols & Integrity Layers

Date: 14/12/2025

Status: Draft / Proposal

Context: Addressing the "Smart Lie" and "Reasoning Drift" risks identified in 2026-era LLMs.

17. The Content Integrity Protocol ("The Immutable Block")

Problem: The Local SLM (Gemma 2B) is tasked with rewriting Cloud outputs into the "Mate" persona. Small models prone to "Instruction Drift" may accidentally alter hard facts (e.g., changing a negative sign or dropping a decimal) while trying to use slang.

Solution: The "Glass Box" UI Architecture.

Mechanism: The Cloud LLM output is structured as JSON with separated fields:

{ "core_truth": "x = -5", "explanation_raw": "Because dividing by a negative flips the sign." }

Rule: The Local SLM has Read-Only access to "core_truth." It cannot rewrite, edit, or "vibe check" this field.

Rendering: The UI renders "core_truth" directly from the JSON into a distinct, high-contrast UI block (The Glass Box).

The Local SLM generates the chat bubble *around* this box: "Check out the answer in the box below, mate!"

Result: The "Vibe" is generated locally, but the "Fact" is piped directly from the Cloud, eliminating the "Blind Toddler" corruption risk.

18. The Logic Verification Layer ("Code-as-Truth")

Problem: "System 2" reasoning models (Gemini 3.0) can hallucinate plausible-sounding but mathematically false derivations ("Smart Lies").

Solution: Mandatory Code Execution (Sandboxed Python).

Workflow: The Cloud Model is forbidden from calculating math via token prediction (internal monologue).

Requirement: It must write and execute a Python script to solve the equation.

Comparison: The Model compares its *intended* answer with the *executed* Python output.

If they differ, it triggers a self-correction loop before sending data to the user.

Verdict: We trust the Python Interpreter (Deterministic), not the LLM (Probabilistic), for the final numerical value.

19. The Syllabus Scope Guardrail ("The Bouncer")

Problem: High-intelligence models often default to university-level methods (e.g., L'Hôpital's rule) which may be valid math but invalid for HSC marking rubrics.

Solution: Negative Constraint RAG.

Implementation: The System Prompt includes a "Negative Knowledge Base"—a list of methods *specifically excluded* from the NSW Year 12 Syllabus.

Check Step: Before outputting a solution, the model runs a "Scope Check": "Does this solution utilize [Forbidden_Method]? If yes, regenerate using [Allowed_Method]."

Example: "L'Hôpital's rule detected. User is Year 12 NSW. Regenerate using First Principles."

20. The "Dead-Switch" Fallback (Offline Continuity)

Problem: The "Guess-First" pedagogy relies on a Cloud response. If the network fails during the 5-second wait, the user is left hanging.

Solution: The "Cached Tutor" Protocol.

Trigger: If Cloud Latency > 8000ms OR Network Error detected.

Action: The Local SLM abandons the "Guess-First" flow and accesses the Local Vector Store (User Notes).

Fallback Response: "Looks like the cloud brain is sleeping (Network Error). Let's check your own notes while we wait. You wrote this about Integration last week..."

Benefit: Keeps the user in the app and studying, even during total service outage.

General things to add idk

General things to add idk

GOOGLE DOCS SUCK AT EQUATIONS AND PDFS ARE INCOMPATIBLE

- Add basically a Google Doc wrapper that actually takes care of functions, and also can be shared by student-tutor and can be worked on live etc.
- Also add PDF support - a new tab (like the functionality that already exists) but also allows a pdf to be inserted natively to the “Google Doc” and you can write on it and add text and images, just cant edit the content directly… unless i can figure out how to make a pdf into a HTML… which seems hard…
- EQUATIONS ARE SO HARD TO ADD AND THEY DON’T LOOK RIGHT and are tedious and ugh
    - Homework questions + answer space with an instant checker…
    - Question generator based on those in the doc
    - Native Gemini support… or just MAIT assistance

General image search in-window

- Just a simple thing that the SLM does a Google Images search for say diagrams of the solar system or the periodic table or anything that is hard to choose a specific image for, and the student will benefit from a preview and link to diagrams from google images
- Like an <iframe> but actually **Google Custom Search Engine (CSE)** or **Custom Search JSON API** integration to securely and reliably embed image results directly into the front-end UI.

Here are the succinct summaries for your document:

Socratic "Guide-Only" Protocol

- Implementation of a system prompt that strictly prevents the AI from providing direct answers. Instead, it functions as a coach: asking guiding questions ("What do you think the first step is?"), offering hints, or verifying a solution only after the student has attempted it.
- Maybe a separate “answer checker” mode where you send the question and your answer and it says correct or not, then you can continue with the learning mode if it’s wrong or even make more questions like it

"Feynman" Reverse-Tutoring Mode

- A flipped interaction where the AI simulates a novice student (intentionally making common misconceptions) and challenges the user to "teach" the concept back to it. This reinforces mastery through the "learning by teaching" method.

SLSO Accessibility & Inclusion Suite

- A dedicated UI toolbar for differentiation, including an instant "Simplify Text" button (rewriting output to a lower reading level), OpenDyslexic font toggles, and Text-to-Speech integration to support diverse learning needs.

Dynamic view for ultra Alded learning

- Current Gemini has this feature but it feels open-ended. I'll add a feature that uses it to pull up photos and make diagrams or show a desmos graph window or whatever to assist the chat other than just the chat window…

Technical Implementation:

- Trigger detection: Confusion markers in chat (repeated questions, "I don't get it", long pauses)
- Can highlight and circle certain sections to trigger Tri-brain image recog and OCR to explain certain sections on a diagram or on text for a particular window
- Content matching: RAG pulls relevant visual from syllabus-aligned media library
- Contextual framing: SLM adds "Look at where X intersects Y" overlay
- Return to flow: Visual minimizes after comprehension check passes